

# Deep Learning at PlusAI

Xiaoshi Wang

# Deep Learning at PlusAI

## History

## Modeling Procedure at PlusAI

- Problem/ Motivation
- Task Formation
- Data/Label Generation
- Modeling
- Evaluation
- Deployment

## Challenges & Future Plans

# A Little Bit of History

2017:

- 2D obstacle detection
- 2D lane detection

2020:

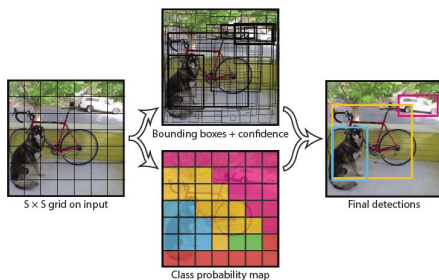
- 2D obstacle detection
- 2D lane detection
- Mono depth estimation
- 3D obstacle detection
- Prediction
- Control (vehicle model)
- ...

# A Little Bit of History


## 2D obstacle and lane detection

Obstacle: YOLOv2 + Multiscale -> YOLOv3 + Multiscale -> CenterNet

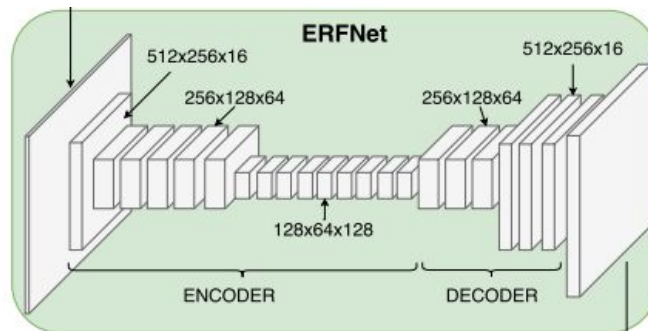
YOLOv3:  
Anchor  
based



CenterNet:  
Anchor free



Lane: ERFNet + techniques



## A Little Bit of History



## Lidar 3D obstacle detection

## Prediction

■ ■ ■

## WIP tasks

- Object Orientation Detection
- Image Scenario Detection
- Keypoint Detection
- Image Object Tracking
- ...



# Modeling Procedure at PlusAI

## Motivation & Task Formation (For new tasks)

### Motivation:

- Things we can improve in current pipeline

### Task Formation:

- Learning/Non learning based.
- Classification



# Modeling Procedure at PlusAI

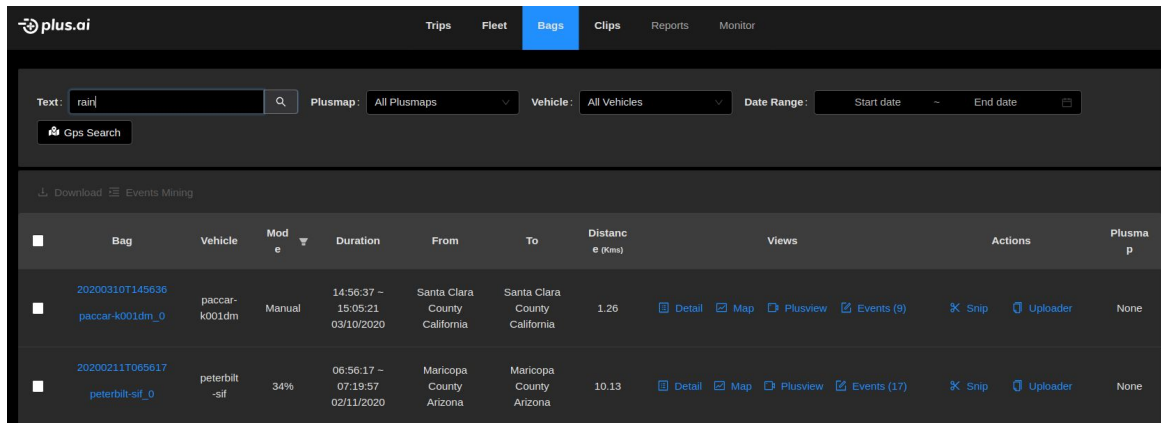
Data Generation : Data gathering

- Data Download (from data1)
- Data Extraction
  - /front\_left\_camera/image\_color/compressed
  - /unified/velodyne\_points
  - /navsat/odom
  - ...

# Modeling Procedure at PlusAI

## Data Generation : Data gathering

- Data Download & Extraction By Events



The screenshot shows the PlusAI web interface. At the top, there's a navigation bar with tabs: Trips, Fleet, Bags (selected), Clips, Reports, and Monitor. Below the navigation bar, there's a search section with a text input containing 'rain', a search button, and filters for Plusmap (All Plusmaps), Vehicle (All Vehicles), and Date Range (Start date to End date). A 'Gps Search' button is also present. Below the search section, there's a table with columns: Bag, Vehicle, Mode, Duration, From, To, Distance (Kms), Views, Actions, and Plusmap. The table contains two rows of data.

Bag	Vehicle	Mode	Duration	From	To	Distance (Kms)	Views	Actions	Plusmap
20200310T145636 paccar-k001dm_0	paccar-k001dm	Manual	14:56:37 ~ 15:05:21 03/10/2020	Santa Clara County California	Santa Clara County California	1.26	<a href="#">Detail</a> <a href="#">Map</a> <a href="#">Plusview</a> <a href="#">Events (9)</a>	<a href="#">Snip</a> <a href="#">Uploader</a>	None
20200211T065617 peterbilt-sif_0	peterbilt-sif	34%	06:56:17 ~ 07:19:57 02/11/2020	Maricopa County Arizona	Maricopa County Arizona	10.13	<a href="#">Detail</a> <a href="#">Map</a> <a href="#">Plusview</a> <a href="#">Events (17)</a>	<a href="#">Snip</a> <a href="#">Uploader</a>	None

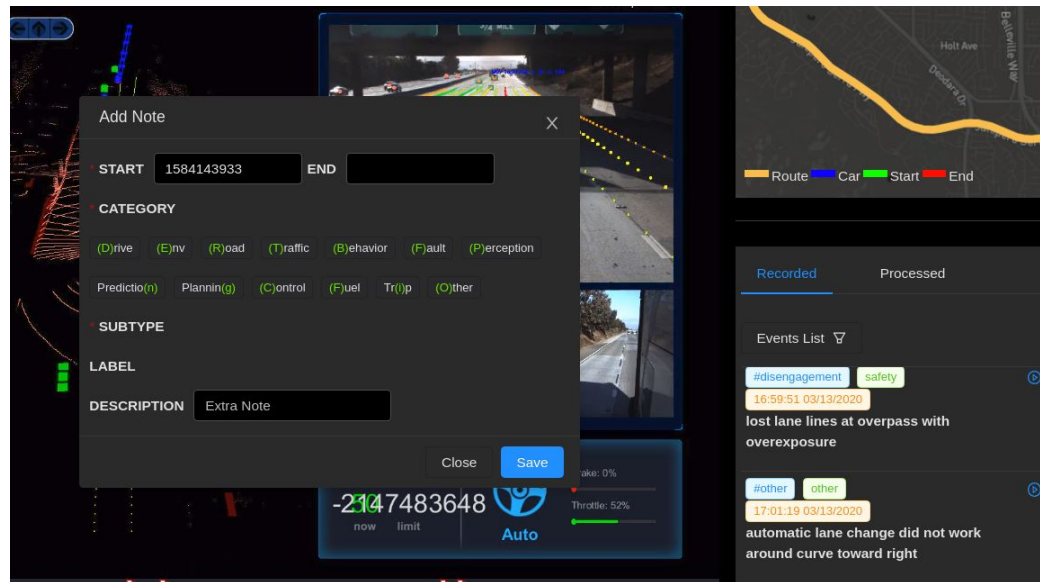
- More Automatic in Future:
  - Advanced Event Processing and Offline Data Mining



# Modeling Procedure at PlusAI

## Data Generation:

- Manual Annotation
  - Taxonomy
- Event Processing
  - Automatically detect and process events in Bags.



- [https://github.com/PlusAI/event\\_processing](https://github.com/PlusAI/event_processing)

# Modeling Procedure at PlusAI

## Data Generation: Labeling

## Procedure

- Data Uploading
- Task Generation
  - Data position
  - Guideline
  - Deadline
- Task Tracking
- Data Verification
  - Labeling team
  - Engineer


```
aws s3 cp data.tar s3://labeling/hour/ --endpoint-url=http://172.16.0.3 --profile suzhou
```

The screenshot displays a Kanban board for 'regular task labeling'. The board is organized into four main columns: 'task submission', 'pre-labeling', 'labeling', and 'internal checking'. Each column contains several task cards. The 'task submission' column has 5 cards, 'pre-labeling' has 3, 'labeling' has 7, and 'internal checking' has 0. Each card typically includes a title, a status (e.g., '2d\_object', 'urgent'), a deadline, and a person icon. The board also features a search bar at the top and a 'see 12 more' link.

Column	Card Title	Status	Deadline
task submission	2020-01-10-to-2020-01-21-selected5	2d_object	12/27
	2020-01-10-to-2020-01-21-selected6	2d_object	12/27
	2020-01-10-to-2020-01-21-selected8	2d_object	12/29
	2020-01-10-to-2020-01-21-selected7	2d_object	12/25
	Image Selection and 2D bounding box Labeling (20190814)	2d_object	12/23
pre-labeling	Image Selection and 2D bounding box Labeling (20190812_2)	2d_object urgent	12/27
	Image Selection and 2D bounding box Labeling (20190813)	2d_object urgent	12/29
	Image Selection and 2D bounding box Labeling (20191031)	2d_object urgent	12/25
labeling	lane detection amazon run/LOL relabel using v2 guideline	lane urgent	4/28
	2020-01-pacaar-k001-undistorted5	2d_object urgent	
	2020-01-pacaar-k001-undistorted4	2d_object urgent	
	2020-01-10-to-2020-01-21-selected4	2d_object	
	2020-01-10-to-2020-01-21-selected3	2d_object	
	2020-02-lewis-selected4	2d_object urgent	
internal checking			

# Modeling Procedure at PlusAI

## Data Generation: Data Verification (QA)




# Modeling Procedure at PlusAI


## Data Generation: Dataset Split, Benchmark Dataset

- Training, Validation, Test (Benchmark)
- Test set is annotated with scenarios

Weather



Exposure



Traffic Status



# Modeling Procedure at PlusAI

## Data Generation: Existing Datasets

### Large Datasets:

- 2D obstacle detection: 100000 frames
- 2D lane detection: 60000 frames
- 3D lidar&image tracking: 50000 frames
- 2D ground touching points (side view): 30000 frames

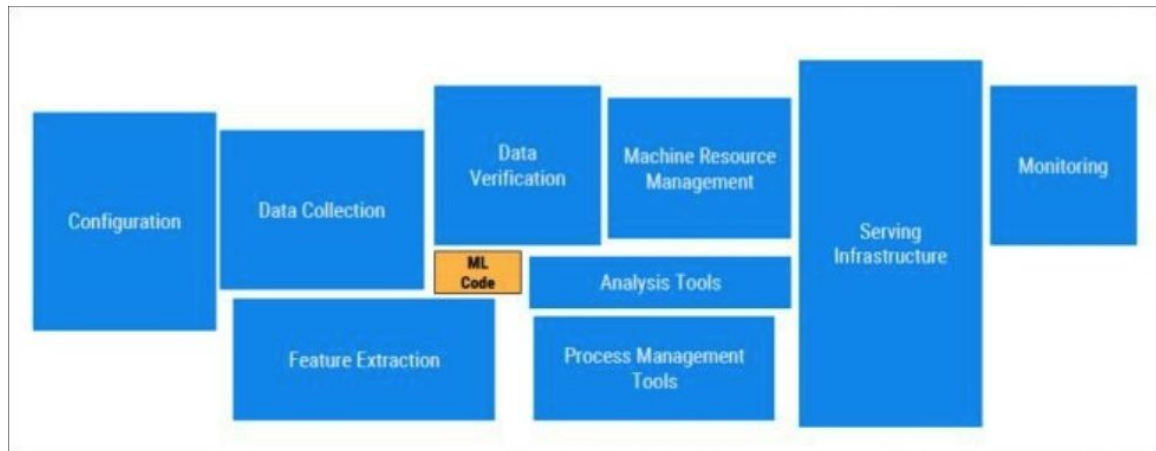
### Small Datasets:

- 2D Key point data (2000-4000 frames in May)
- Semantic & instance segmentation data (200-500 in May)

# Modeling Procedure at PlusAI

## Modeling -- Coding

- Data pipeline
- Model Architecture
- Evaluation Metrics
- Training Script



# Modeling Procedure at PlusAI

## Modeling -- Training

Hardware: Suzhou Gpu Clusters, AWS

```
xiaoshiw@gpu6:~$ nvidia-smi
Wed Apr 29 11:26:43 2020

+-----+
| NVIDIA-SMI 410.129      Driver Version: 410.129      CUDA V
+-----+-----+
| GPU   Name                Persistence-M| Bus-Id        Disp.A | Vol
+-----+-----+
| Fan   Temp   Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-
+-----+-----+
|  0    GeForce RTX 208...    On      | 00000000:1A:00:0 Off |
+-----+-----+
| 29%   49C    P2      93W / 250W | 10947MiB / 10989MiB |
+-----+-----+
```

## Training Procedure

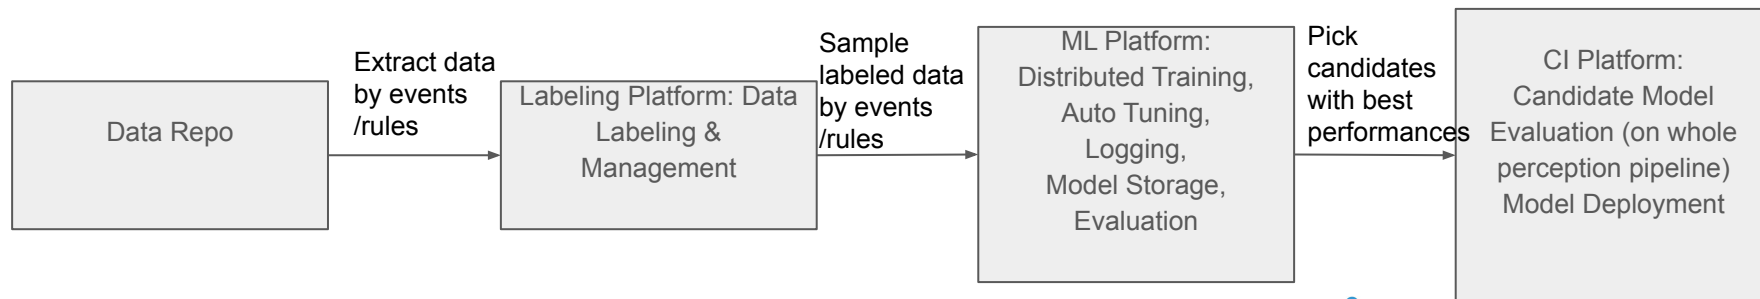
- Structure & capacity
- Loss
- Scale
- Tricks (hard negative sampling etc.)
- Hyperparameter tuning

## Model

Name	Status	Performance	Note
Anchor regeneration	Finished for small scale model; In progress for large scale	Positive	The improved small scale model has been merged.
RefineNet	Finished	Negative	
RetinaNet	Implemented		Planned for detailed examination in future
Yolo v3 (Darknet43) + multiscale output	Implemented	Positive for faraway objects  Testing in progress	
Yolo v3 (Darknet26) + multiscale output	Implemented	Positive for faraway objects  Testing in progress	
Network random search	Implemented v1		Planned for detailed examination in future
Focal loss	Finished	Positive	

# Modeling Procedure at PlusAI

Modeling -- Training ML Platform (in recent future)



kubernetes



Kubeflow





# Modeling Procedure at PlusAI

## Evaluation Metrics:

- Thorough understanding of model
- Product orientated

## Obstacle Detection:

- Goal: Large objects, nearby objects, ego lane objects, pedestrian, robustness
- Metrics Design:
  - metrics based on obstacle size,
  - truncated obstacles,
  - occluded obstacles
  - front obstacle,
  - ego lane obstacles,
  - Robustness of objects in continuous frames

# Modeling Procedure at PlusAI

## Evaluation: PlusAI vs KITTI

	PlusAI	KITTI
Metrics	mAP, Precision, Recall, F scores, Confusion Matrix	mAP
Scenarios	Large object (bucketized); Truncated; Occluded; Ego Lane Object (fp); Front Vehicle (recall); Robustness;	Truncated + Occluded
Deployment	Class & size weighted F1 score with precision & recall restrictions	

```
"box_size": "0.003-0.004",
"result": {
  "num_images": 4757,
  "prob_threshold": 0.5,
  "eval_threshold": 0.5,
  "mAPs": {
    "car": 0.9567511980608188,
    "bus": 0.28188229696109174,
    "truck": 0.762556853647838,
    "moto": 0.004201659663970588,
    "pedestrian": 0.9473679224379354,
    "bike": 0.8333291666874998,
    "traffic_cone": 0.8461531952667729,
    "barricade": 0.368865533361575
  },
  "num_vehicles": {
    "car": 1959,
    "bus": 34,
    "truck": 264,
    "moto": 2,
    "pedestrian": 19,
    "bike": 2,
    "traffic_cone": 13,
    "barricade": 332
  },
  "num_front_vehicles": 180,
  "front_car_recalls": {
    "0.8": 0.89444444439475309,
    "0.85": 0.89444444439475309,
    "0.9": 0.89444444439475309,
    "0.95": 0.89444444439475309
  },
  "all_recalls": {
    "0.8": 0.8205714285401687,
    "0.85": 0.814095238064225,
    "0.9": 0.8076190475882812,
    "0.95": 0.7912380952079529
  },
  "F_scores": {
    "0.5": 0.9185991387898435,
    "1.0": 0.8420824245937985
  },
  "precision": {
    "car": 0.9858916478499669,
    "bus": 0.999999995,
    "truck": 0.8989361701649502,
    "pedestrian": 0.9999999933333333,
    "traffic_cone": 0.9999999985714286,
    "barricade": 0.9999999900000002,
    "all": 0.9778337531436886
  },
  "recall": {
    "car": 0.8917815211797255,
    "bus": 0.05882352939446367,
    "truck": 0.640151515127267,
    "pedestrian": 0.7894736837950138,
    "traffic_cone": 0.5384615380473372,
    "barricade": 0.0030120481926803598,
    "all": 0.7394285714257546
  },
  "ego_lane_fp/total_fp": 0.6454116572810145,
  "ego_lane_fp/total_images": 0.891528274122346,
  "confusion_matrices on size-matched detection
  (or gt, col for det)": [
    [
      1747, 0, 14, 0, 0, 0, 0,
      3, 2, 1, 0, 0, 0, 0,
      6, 0, 169, 0, 0, 0, 0,
      2, 0, 0, 0, 0, 0, 0,
      0, 0, 0, 0, 15, 0, 0,
      0, 0, 0, 0, 0, 0, 0
    ]
  ]
}
```

# Modeling Procedure at PlusAI

## Evaluation -- Analysis

Why the new network/technique performs better

### Why is CenterNet better than Yolov3?

Hypotheses	Possible effects	How to validate?
<u>CenterNet</u> has a better backbone.	++ small objects ++ large objects	Replace the darknet backbone of Yolov1 with DLA
<u>CenterNet</u> uses soft labels for the locations while Yolov3 uses hard labels.	+ small objects + large objects	Hard to validate
The regression loss function of CenterNet cares more for large objects.	-- small objects ++ large objects	Use a similar loss function for Yolov3
Small bounding boxes are not filtered out during training of <u>CenterNet</u> .	++ small objects - large objects	Train Yolov3 without filtering out small objects. <b>Validated</b>

# Deployment

## Runtime API

## Model Speedup

## PR Submission

```
class RObjectDetector : public ObjectDetector {
public:
    typedef std::shared_ptr<RObjectDetector> Ptr;

    // async function. will submit jobs to GPU and return immediately.
    void DetectAsync(const cv::cuda::GpuMat& img) override;
    // this function will block until GPU job finished.
    void GetDetection(std::vector<models::GenericDetection>& detections) override;
    // this function will block until GPU job finished.
    void GetDetection(std::vector<GenericDetection3D>& detections_3d) override;
    // this function will block until GPU job finished.
    void GetFeatures(cv::Mat& features) override;
    bool IsDetectionFinished() override;

    RObjectDetector(const StereoDetectorParam& conf,
                    const boost::filesystem::path& file_root_path,
                    bool get_image_features = false);

private:
    std::unique_ptr<TensorrtDetector> _detector;
    const StereoDetectorParam _conf;
    int _gpu_id;
};
```

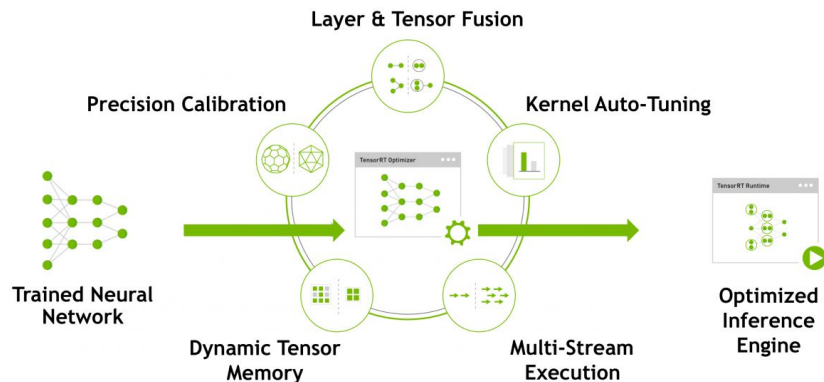
Metric	Equation	Overall	close_front_ego_ego_lane x: (2, 50) ,  y : (0, 2)	close_front_neighbor_lane x: (2, 50) ,  y : (2, 6)	medium_front_ego_lane x: (50, 100) ,  y : (0, 2)	medium_front_neighbor_lane x: (50, 100) ,  y : (2, 6)	far_front_ego_lane x: (100, 150) ,  y : (0, 2)
Recall	$TP / (TP + FN)$	0.591 -> 0.586	1.000 -> 1.000	0.996 -> 0.995	0.988 -> 0.981	0.961 -> 0.956	0.956 -> 0.959
Precision	$TP / (TP + FP)$	0.750 -> 0.741	1.000 -> 1.000	0.976 -> 0.974	0.932 -> 0.949	0.964 -> 0.968	0.885 -> 0.888
Matching Rate	$1 - MM / TP$	1.000 -> 1.000	1.000 -> 1.000	1.000 -> 1.000	1.000 -> 1.000	1.000 -> 1.000	1.000 -> 1.000
Multi Object Tracking Accuracy	$1 - (MM + FP + FN) / (TP + FN)$	0.394 -> 0.381	1.000 -> 1.000	0.971 -> 0.969	0.915 -> 0.928	0.925 -> 0.925	0.832 -> 0.838
3d cuboid iou	$mean(iou3d(tp\_has\_lidar, gt\_has\_lidar))$	0.193 -> 0.194	0.395 -> 0.391	0.492 -> 0.492	0.193 -> 0.193	0.161 -> 0.162	0.024 -> 0.026
2d box iou	$mean(iou2d(tp\_has\_camera, gt\_has\_camera))$	0.692 -> 0.691	0.866 -> 0.865	0.828 -> 0.828	0.787 -> 0.780	0.782 -> 0.777	0.655 -> 0.656

# Deployment

Model Speedup: TensorRT, Cuda post processing

## TensorRT

- Layer / Tensor Fusion
- Auto-Tuning
- Precision Calibration
- Multi-Stream Execution
- Dynamic Tensor Memory



Offline: TensorFlow -> UFF, PyTorch -> ONNX

Online: UFF / ONNX -> TensorRT Engine

## Cuda post processing

- Run lots of tasks in parallel

# Challenges & Future Plans

## Data Filtering

- event processing, online & offline

## Model Latency

- multitask learning

## Model Tuning

- distributed training, AutoML