

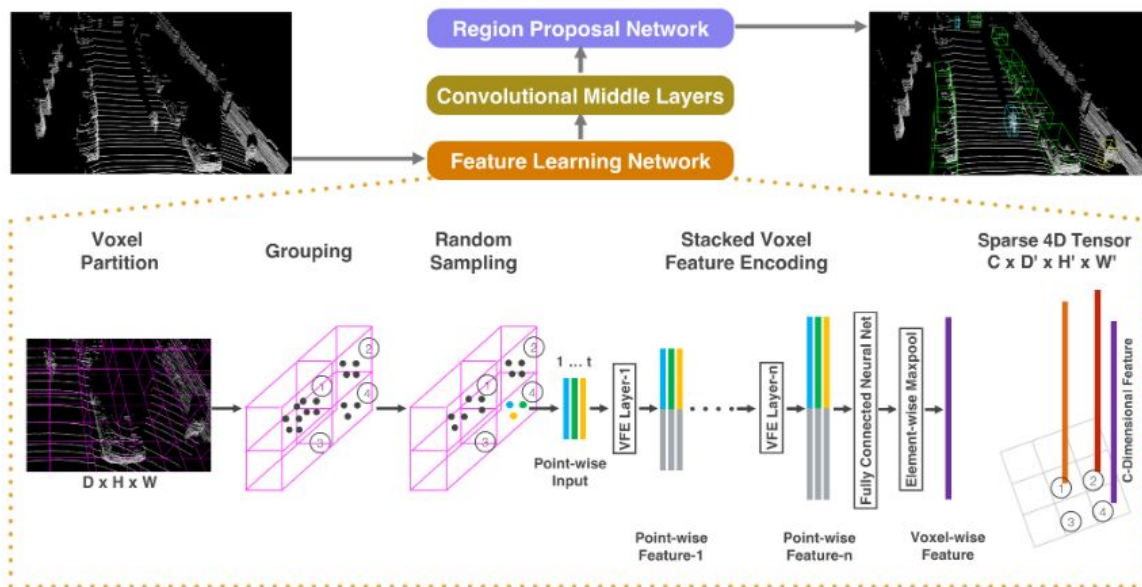
3D obstacle detection

Pointpillars

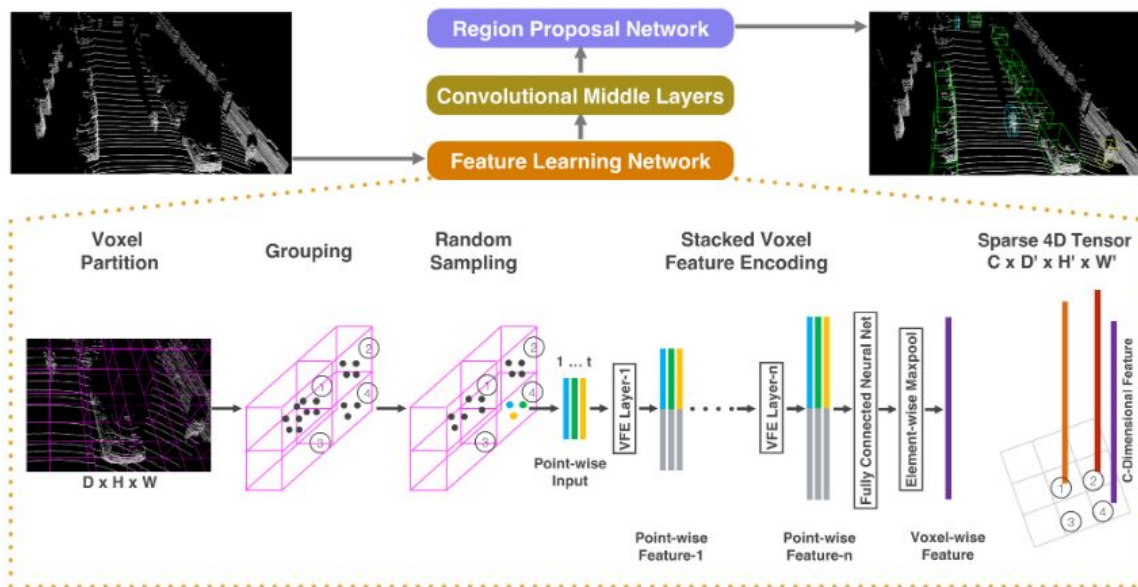
Introduction

- LiDAR
 - Reliable depth estimation
 - Accurate localization
- Very sparse, variable density
- Region Proposal Network
 - Highly optimized for efficient object detection
 - Needs input in the form of images !!
- Backbone needs to extract features reliably in pseudo-image form

VoxelNet




VoxelNet



- 3D convolutions in the middle layers make it too slow !

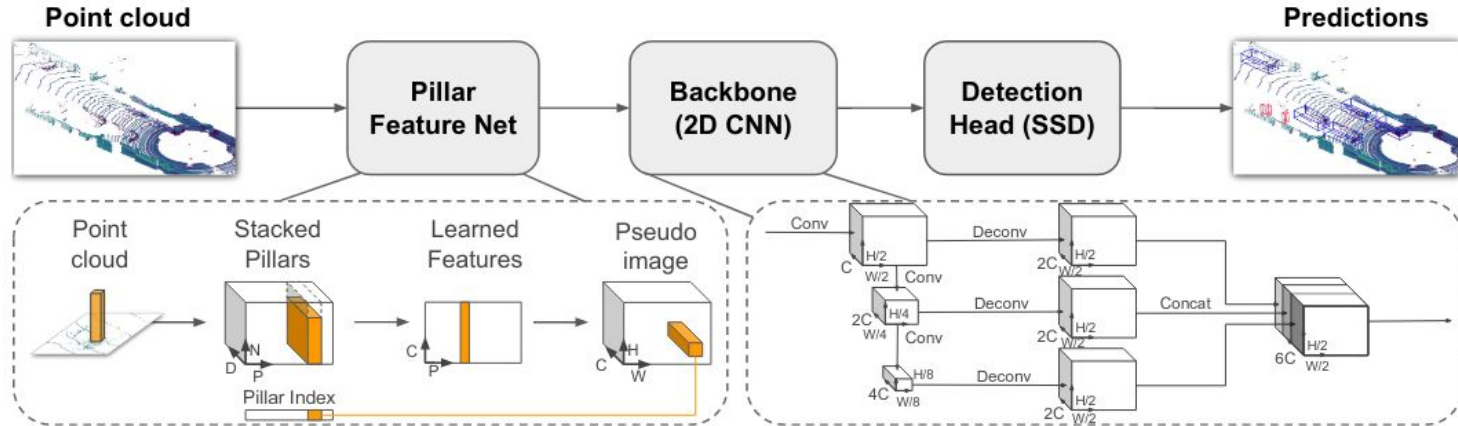
PointPillars

- Instead of voxels, learn features on pillars
- P: Number of non-empty pillars
- N: Max number of points per pillar
- D: Augmented lidar point dimension (=9)
- No 3D convolutions !!
- Very fast! (4.4 Hz -> 100 Hz)



PointNet -> Pillar-wise Max

Pointpillars architecture




Runtime deployment: model and CUDA kernels

We use onnx as the model parser and use tensorRT for inference

Sadly not all operations during inference are onnx compatible:

Pillarization, Pseudo Image generation, Postprocessing need their own CUDA kernels



All kernels are present: *perception/obstacle_detection/src/lidar/point_pillars/*.cu*

Runtime deployment: model config

https://github.com/PlusAI/drive/blob/master/perception/config/fusion_tracker/fusion_tracker.prototxt.paccar

```
-----  
# lidar_nn_param {  
# Inference general settings  
  gpu_device: 0  
# TensorRT engine settings  
  engine_pfe {  
    # This engine should use FP32 or FP16  
    uff_file: "/opt/plusai/var/models/obstacle_detection/pointpillars_pfe_us_v1.2.onnx"  
    data_type: FP16  
    input_channels: 10  
    input_height: 40960  
    input_width: 32  
    input_nodes: "input"  
    detection_output_nodes: "output"  
  }  
  engine_rpn {  
    # This engine can be switched among FP32, FP16 and INT8  
    uff_file: "/opt/plusai/var/models/obstacle_detection/pointpillars_rpn_us_v1.2.onnx"  
    data_type: FP16  
    input_channels: 64  
    input_height: 80  
    input_width: 400  
    input_nodes: "input"  
    detection_output_nodes: "output"  
  }  
}  
# Model information  
  num_class: 1  
  batch_size: 1  
  range {  
    min_x: -14.0  
    min_y: -25.6  
    min_z: 2.0  
    max_x: 50.0  
    max_y: 25.6  
    max_z: 6.0  
  }  
# Anchor generation  
  anchor_stride: 2  
  num_anchors: 40960  
  num_anchor_rotation: 2  
  anchor_rot_angles: 0  
  anchor_rot_angles: 1.5708  
  num_dir_bins: 2  
  dir_offset: 0.78539  
# CAR anchor info  
  anchor_sizes: 4.63 # length  
  anchor_sizes: 1.97 # width  
  anchor_sizes: 1.74 # height  
  anchor_center_heights_z: 0.47  
# Pillarization  
  pillar_size {  
    x: 0.2  
    y: 0.2  
    z: 0.0  
  }  
  max_num_pillars: 24000  
  max_num_points_per_pillar: 32  
  zero_out_intensity: false  
  num_gathered_point_features: 10  
# Model inference  
  num_pfe_output_features: 64  
# Postprocess  
  prob_threshold: 0.3  
  nms_threshold: 0.01  
  max_num_boxes_for_nms: 1024  
}
```


Losses

- Localization loss

$$\mathcal{L}_{loc} = \sum_{b \in (x, y, z, w, l, h, \theta)} \text{SmoothL1}(\Delta b)$$

$$\begin{aligned}\Delta x &= \frac{x^{gt} - x^a}{d^a}, \Delta y = \frac{y^{gt} - y^a}{d^a}, \Delta z = \frac{z^{gt} - z^a}{h^a} \\ \Delta w &= \log \frac{w^{gt}}{w^a}, \Delta l = \log \frac{l^{gt}}{l^a}, \Delta h = \log \frac{h^{gt}}{h^a} \\ \Delta \theta &= \sin(\theta^{gt} - \theta^a),\end{aligned}$$

- Bin Classification Loss:
- Object Classification Loss

Softmax

$$\mathcal{L}_{cls} = -\alpha_a (1 - p^a)^\gamma \log p^a,$$

Total Loss => Weighted sum

Quantitative results on Plus Dataset



- As of 05/01/2022, our 3D dataset only has 1 Class (Car/Truck) labeled

mAP	Precision	Recall
88.91	90.55	90.01

- Avg errors in prediction metrics

X_error (m)	Y_error (m)	L_error (m)	W_error (m)	Yaw_error (rad)
0.21	0.13	0.30	0.11	0.13

Qualitative Results on Plus Dataset



Tools and Repos

LidarDet repo

Repo for 3D obstacle detection training

<https://github.com/PlusAI/LidarDet>

Models repo

Repo containing trained models which are used at runtime

<https://github.com/PlusAI/models>

Labeling guideline for 3D obstacle detection

[https://github.com/divimund/tools/blob/master/labeling/guidelines/3d_object_tracking_labeling_guideline.m
d](https://github.com/divimund/tools/blob/master/labeling/guidelines/3d_object_tracking_labeling_guideline.md)

Labeled data available at


s3://labeling/benchmark/obstacle_tracking/data/

aws s3 --endpoint-url=http://172.16.0.3 ls s3://labeling/benchmark/obstacle_tracking/data/ --recursive --human-readable --summarize --page-size=100000

3D detection debugging

Run Pointcloud Perception Util

```
python ./perception/obstacle_detection/tools/run_pointcloud_perception.py \  
-- bag {bag_path}
```



Run Perception Stack

Unified perception simulator

```
python ./perception/simulation/scripts/run_unified_simulator.py <bag_path>
```